

[WSE 380] Technical Foundations of a Startup

Secure Web Application (WordPress) Management

Instructors: Johnny So and Professor Nick Nikiforakis

Contents

1	Session 0: SSH Connections	3
1.1	Course Materials	3
1.2	Installing an SSH Client	3
1.3	Connecting to a Remote Machine	3
1.3.1	SSH Keys	3
1.3.2	SSH Config	4
2	Session 1: The Command Line & Prerequisites	5
2.1	Linux Command Line Basics	5
2.1.1	Command Line Anatomy	5
2.1.2	The Working Directory	5
2.1.3	Moving Around the File System	6
2.1.4	Creating, Reading, Moving, Copying, and Deleting Files and Directories	6
2.1.5	Editing Files	7
2.1.6	Downloading Programs	7
2.1.7	Downloading Code Using Git	7
2.1.8	Running Programs	7
2.1.9	Sudo	8
2.1.10	Tab Completion	8
2.2	Setup and Administration of a Server	8
2.2.1	SSH	9
2.2.2	UFW (Uncomplicated Firewall) [4]	9
2.2.3	Fail2Ban [4]	10
2.3	Installing Docker and Docker Compose	10
3	Session 2: Web Application Setup	12
3.1	Web Application Infrastructure	12
3.2	Docker Compose Services	12
3.2.1	Environment Variables	12
3.3	WordPress Installation via the Web Interface	13
3.4	DNS	13
3.4.1	Domain Registration	13
3.4.2	DNS Name Servers	14
3.4.3	Updating the Nginx/WordPress Configuration	14
3.5	TLS Certificates	14
3.5.1	Certificate Request	14
3.5.2	Updating Docker Compose	15
4	Session 3: WordPress Customization	16
4.1	Content	16
4.1.1	Categories	16
4.1.2	Tags	16
4.1.3	Posts	16
4.2	Styling and Themes	18
4.3	Plugins	18
4.3.1	Security - Wordfence	18
4.4	Free Time: Site Customization	19
5	Session 4: Website/Server Monitoring	20

5.1	Resource Usage	20
5.2	Uptime	20
5.3	More Plugins	20
5.3.1	Analytics - Google Site Kit	21
5.3.2	Performance - LiteSpeed Cache (Optional)	21
5.3.3	Forms - WPForms (Optional)	21
5.3.4	Forms - WP Mail SMTP (Optional)	21
5.4	Free Time: Site Customization	22
6	Session 5: Server Traces & Site Showcase & Review	23
6.1	Understanding Logs	23
6.1.1	Server Logs	23
6.1.2	Application Logs.	23
6.2	Site Showcase	24
6.3	Conceptual Review	24

1 Session 0: SSH Connections

Welcome to WSE 380: Technical Foundations of a Startup! Today, we will be setting up your devices with an SSH client to connect to a virtual machine that we will be running experiments on.

1.1 Course Materials

All of the course materials are stored at the following GitHub repository: <https://github.com/jso123450/wse380-webapp.git>.

1.2 Installing an SSH Client

SSH (Secure SHell) is a network protocol that enables a computer to establish a secure connection to a remote computer that is running an SSH server. To connect over SSH, we need to make sure we have an SSH client on our devices. New versions of Windows (10 and later) and MacOS come with built-in SSH clients, so there is no need to install a separate client; Windows (10 and later) users can use the Windows Terminal application whereas MacOS users can use the Terminal application. However, if you are running an older version of Windows or if you would like to install a separate client, you can install MobaXTerm (recommended) or PuTTY.

1.3 Connecting to a Remote Machine

The default syntax for the command to establish an SSH connection is `ssh -p 22 user@address`, where

- `ssh` is the name of the program,
- `-p 22` is a flag `-p` with a value of 22 (for SSH, this flag indicates the port),
- `user` is the username of the remote account, and
- `address` is an address to the remote machine. This can be a domain name (e.g., `google.com`) or an IP address (e.g., `1.2.3.4`) — in our case, we will be using IP addresses to connect to our remote machines.

As an example, if you are trying to establish an SSH connection to a remote machine with IP address `1.2.3.4` on port 5678 with a user account of `foobar`, you would execute the following command in your terminal:

```
1 user@local-machine:~$ ssh -p 5678 foobar@1.2.3.4
```

If you are authenticating with a username and password, the SSH client should then prompt you to enter a password:

```
1 foobar@1.2.3.4's password:
```

1.3.1 SSH Keys

We will be connecting over SSH using a public and private key pair to eliminate the use of a password. You can read more about them on this DigitalOcean tutorial or on this GitHub documentation. Let us open a terminal and enter the following:

```
1 user@laptop:~$ ssh-keygen -t ed25519 -C "your_email@example.com"
```

For the following prompts that ask about the file location and key passphrase, you can hit `Enter` without specifying anything to accept the default values (the file location will be `~/.ssh/id_ed25519` and there will be no passphrase).

After you have done this, please do the following:

```
1 user@laptop:~$ cat ~/.ssh/id_ed25519.pub
2 ssh-ed25519 <random_string> <your_email@example.com>
```

Please copy the output and paste it in a reply to the email with your login credentials.

1.3.2 SSH Config

To facilitate ease of working with remote machines, we will configure a SSH configuration file and assign a name to your own machine. The default SSH configuration file is located at “`~/.ssh/config`”, where `~` is the home directory of your user. Let us open a terminal and enter the following:

```
1 user@laptop:~$ nano ~/.ssh/config
```

This should open a new (or edit the existing) file at this location using the basic nano file editor. Copy and paste the following into the terminal:

```
1 Host wse380
2     HostName <ip address from email>
3     Port <port from email>
4     User <username from email>
```

Now press *Ctrl+x*, *y*, *Enter* (hold down the *Ctrl* key and then press *x*, let go of *Ctrl*, then press *y* and *Enter*). This should prompt nano to close and save the file. Now, you should be able to type `ssh wse380` instead of `ssh -p <port> <username>@<ip address>`. Please verify that this works successfully at the beginning of our next session!

2 Session 1: The Command Line & Prerequisites

Welcome to the first hands on session of WSE 380: Technical Foundations of a Startup! Today, we are going to start by talking about the Linux command line and securing our servers.

SSH Connections

Before we begin, make sure that you are able to establish an SSH connection into your provided machines!

2.1 Linux Command Line Basics

Before we dive into deploying our own web applications, we will learn some basics of the Linux command line to gain more familiarity with server management.

2.1.1 Command Line Anatomy

The most basic access that we can have to a computer is through the command line, which is sometimes also called a terminal or a shell. Users are created on machines with certain permissions, and these users can login to the machine via a terminal, and enter commands for the machine one line at a time. Different terminals may present minute differences in the appearance of a command line, but they all share a “command prompt” that looks something like this:

```
1 <username>@<machine_name>:<working_directory>$
```

This “command prompt” comprises:

- `username`: the name of the logged-in user,
- `machine_name`: the name of the machine,
- `working_directory`: the current directory (further explained below), and
- `$`: the symbol which separates the prompt from the command that you enter.

2.1.2 The Working Directory

When using any command line environment, it is up to the user to keep a mental map of where in the file system they are currently located (the “working directory”). Luckily, there are two easy commands to help us figure out where we are and what files and sub directories are currently around us.

- **pwd** - Lists the absolute path of the current directory

```
1 user@wse380:~$ pwd
2 /home/user/
```

- **ls** - Lists the contents of a directory. Specifying no path will display the contents of the current directory.
 - `ls -l`: Displays additional information about each file such as permissions and file size
 - `ls -a`: Displays all files in a directory, including hidden ones

```
1 user@wse380:~$ ls -l -a
2 total 24
3 drwxr-xr-x 2 wse380 wse380 4096 Mar 11 13:00 .
4 drwxr-xr-x 9 root root 4096 Mar 11 13:00 ..
5 -rw----- 1 wse380 wse380 5 Mar 11 13:00 .bash_history
6 -rw-r--r-- 1 wse380 wse380 220 Mar 11 13:00 .bash_logout
7 -rw-r--r-- 1 wse380 wse380 3771 Mar 11 13:00 .bashrc
8 -rw-r--r-- 1 wse380 wse380 807 Mar 11 13:00 .profile
```

Above, we used `ls` to list the contents of the current directory while also specifying the `-l` and `-a` flags to show us all files as well as details on each file. Here, from left to right on each row, we can see the permissions on each file and directory, the user that owns it, the user group that has access to it, its size in bytes, the date and time it was last modified, and its name.

2.1.3 Moving Around the File System

One of the most important symbols in the Linux command line is the “.”, or period symbol. In the context of the file system, “.” represents the directory the user is currently in, and “..” represents the parent directory of the current directory. For the most part, when attempting to navigate the file system or interact with files, you will use file paths relative to your current directory. These are called *relative paths*. For example, if a user was trying to list the contents of the current directory’s parent, they would use the following command:

```
1 user@wse380:~$ ls ../
```

The notion of relative file paths becomes vital when attempting to navigate the file system. To accomplish this, we make use of the “change directory” or `cd` command.

- **cd** - Change the current directory to the one specified. Supplying no arguments navigates to the user’s home directory (the symbol `~` represents the home directory).

```
1 user@wse380:~$ ls
2 wse380-webapp/
3 user@wse380:~$ cd wse380-webapp/
4 user@wse380:~/wse380-webapp$ pwd
5 /home/user/wse380-webapp
6 user@wse380:~/wse380-webapp$ cd ../
```

Above, we issued a series of commands to help visualize our movement around the file system. We first used `ls` to view the files and sub directories we had available to us from our current directory. We then used the `cd` command to move to the “wse380-webapp” sub directory. A call to the `pwd` command shows that our new current directory was now the directory located at “/home/user/wse380-webapp”. Finally, we used the `cd` command with the relative path “..” to move back to the parent directory, the original directory we started at. An example of an *absolute path* would be the output of the earlier `pwd` command: “/home/user/wse380-webapp”. We could have, for example, executed `cd /home/user/wse380-webapp` from anywhere in the system instead of needing to execute `cd wse380-webapp` from our home directory. This is powerful because absolute paths enable us to *refer to any part of the file system, regardless of our current working directory*.

2.1.4 Creating, Reading, Moving, Copying, and Deleting Files and Directories

- **touch** - Create a file with the specified name.

```
1 user@wse380:~$ touch example.txt
```

- **cat** - Print a file’s contents to the terminal window.

```
1 user@wse380:~$ cat example.txt
```

- **mkdir** - Create a new directory with the specified name.

```
1 user@wse380:~$ mkdir exampleFiles
```

- **mv** - Move a file from one location in the file system to another. Also the method used to rename files from the command line.

```
1 user@wse380:~$ mv example.txt exampleFiles/example.txt
```

- **cp** - Copy a file from one location in the file system to another while keeping the original file untouched. When copying a directory, the user must supply the “-R” argument.

```
1 user@wse380:~$ cp example.txt exampleFiles/example.txt
```

- **rm** - Remove a file from the file system. To remove a directory, the user must include the “-r” argument to *recursively* remove all sub directories and files. **IMPORTANT: Removing files with this command is permanent! This command is the equivalent of using your graphical interface to move a file to the trash and emptying it. Only remove a file or directory you are sure you want permanently deleted!**

```
1 user@wse380:~$ rm -r exampleFiles/
```

2.1.5 Editing Files

You may be used to writing code and editing configuration files using an IDE. While many of the most popular IDEs are available on Linux, it is important to know how to edit files from the Linux command line since you will not be able to use a graphical text editor or IDE to edit files on remote servers.

The most popular command line text editors on Linux include: Vim, Nano, and Emacs. We will be using Nano in this course as it has the easiest learning curve. To get started, run the following command to open a file with Nano:

```
1 user@wse380:~$ nano example.txt
```

Here, *example.txt* can either be a new file or an existing file. You will then be able to enter any new text or edit any existing text in the file. Once you are ready to save your work, press *Ctrl+x* to exit Nano. You will be asked if you'd like to save the changes that you made. If so, press *Y*, otherwise, press *N*. Lastly, you will be asked to provide the name of the file you'd like to write your changes to. By default, this will be the name of the file you entered when opening Nano. Most of the time we'd just press enter here. However, if you'd like to keep the original file the same as when you opened it and write your changes to a new file, enter the name of the new file and press enter.

2.1.6 Downloading Programs

When you downloaded and installed applications on your computer in the past, you most likely did so either through an application marketplace or by downloading them from a vendor's website and running an installer program. On Linux, you can download and install programs using the command line program *apt* or “Advanced Package Tool”. In many ways, *apt* is easier to use than the process you are familiar with. The only command we have to learn is:

```
1 user@wse380:~$ sudo apt install program_name
```

Here, “program_name” is the name of the program you would like to install. In order to install applications using *apt*, we must run it with administrator privileges. We will discuss this in Section 2.1.9.

2.1.7 Downloading Code Using Git

As you may be familiar with from your own programming assignments, *git* allows us to upload our code to the cloud for sharing and collaboration. For this course, we are only interested in how to download code from a remote *git* repository to our machine.

Let us clone the repository with the material for this semester by running:

```
1 user@wse380:~$ git clone https://github.com/jso123450/wse380-webapp.git
```

This repository will contain most, if not all, of the files and lesson plans for this course.

2.1.8 Running Programs

After we download a program or write our own, we want to actually run them. Doing so is easy, although the syntax of the command differs depending on what kind of program it is. If a program was downloaded using *apt*, it can be executed simply by typing the name of the program into the terminal, because the program is downloaded to a directory that our command line will check by default. For example, if we downloaded the Internet browser Firefox using *apt* we could run it like so:

```
1 user@wse380:~$ firefox
```

To see where such programs are downloaded, we can use *which*:

```
1 user@wse380:~$ which firefox
2 /usr/bin/firefox
3 user@wse380:~$ which abc
4 abc not found
```

If we write our own program or download a program from another source like GitHub, we need to address the executable file directly in the terminal using a relative or absolute path, as we talked about earlier. For example, if we downloaded a GitHub repository containing an executable file named *foo*, we could run it by first navigating into the directory containing *foo*, and executing the command:

```
1 user@wse380:~$ cd some-random-folder
2 user@wse380:~/some-random-folder$ ./foo
```

Alternatively, we could have used a different relative path, or an absolute path:

```
1 user@wse380:~$ ./some-random-folder/foo
2 user@wse380:~$ /home/user/wse380/some-random-folder/foo
```

When executing any kind of program, there will most likely be data outputted to the terminal window. There are two standard places where a program writes its output: standard out (*stdout*) for regular program output and standard error (*stderr*) for error and log messages. Sometimes we are interested in saving this output to a file for later processing, especially for long-running programs. To do so, we will use the “>” symbol to tell Linux to redirect the *stdout* of the command to a file of our choosing. Take for instance the following command:

```
1 user@wse380:~$ ./foo > foo_output.log
```

Here, instead of *stdout* of *foo* going to the terminal window, it will now appear in a new file called *foo_output.log* in the current working directory. If we are also interested in the *stderr* of the program, we can run it as follows:

```
1 user@wse380:~$ ./foo > foo_output.log 2>&1
```

Here, we redirect both *stdout* and *stderr* of *foo* to the file *foo_output.log* in the current working directory.

2.1.9 Sudo

So far, we learned basic Linux commands we need to setup and manage our servers. However, some commands require administrator, or *root*, privileges. On Linux machines, the root user has permission to do just about anything it wants, even modify or delete critical operating system files. Due to this immense power, we typically don’t directly sign in with the root user account, but rather permit regular user to execute commands as root as they see fit by explicitly stating their intent to do so. This prevents users from accidentally performing dangerous commands without thinking.

To state our intent to run a command as root, we prepend the keyword *sudo* (“super user do”) before the command we would like to run. For example, if we would like to open a file in Nano with root privileges, we would enter the following command:

```
1 user@wse380:~$ sudo nano example.txt
```

You will sometimes be asked to enter your password when you attempt to run a *sudo* command. Just enter the same password you used to log into your account, and you should be good to go, assuming your user account is permitted to execute *sudo* commands.

2.1.10 Tab Completion

Most shells offer *tab completion* — if you begin to type part of a file and press the tab key, the shell should autocomplete the name for you (assuming there is only one file that matches the prefix you have typed). If the prefix matches multiple files, you can press tab multiple times, and your shell should attempt to list all files that match the prefix and cycle through them. Tab completion will only work for files that are contained in directories that the shell is configured to check, plus the current working directory.

2.2 Setup and Administration of a Server

Now that we have learned the basics of the Linux command line, we are going to put everything together to setup our web app server. As previously discussed, setting up a web application in a secure manner requires several different pieces of software. We will start by configuring access control parameters to the server.

2.2.1 SSH

The Secure Shell Protocol or *SSH* is a networking protocol that allows users to connect to a remote computer over a secure channel. When you connect to a SSH server, you will be presented with a command line to execute the same commands we discussed in part 1. Since compromising a host over SSH allows for access to all programs and files of a particular user, attacks against this protocol are very common. By default, SSH servers listen on port 22, and this is the typical port that attackers choose to probe for vulnerable SSH servers. There are a number of other default ports that attackers commonly scan, and potentially vulnerable services on different ports. We will be configuring a firewall to limit which ports are accessible, and an intrusion prevention software to monitor the SSH port 22, and others, for connection requests and the appropriate enforcement actions. Another technique that can defend against such probes (that we are currently using) is to change the port that the SSH service is running on from the standard value of 22 to an arbitrary one.

2.2.2 UFW (Uncomplicated Firewall) [4]

Firewalls are network security systems that monitor and control incoming and outgoing traffic according to a set of specified rules. They can come in the form of software, or a dedicated physical machine. In this course, we will be using a ubiquitous Linux software firewall called *ufw* (uncomplicated firewall) to setup rules to restrict access to our VMs. In particular, we want to make sure that we are only exposing access to the necessary ports.

Let's install *ufw* with the following:

```
1 sudo apt update
2 sudo apt upgrade
3 sudo apt install ufw
```

Now, let us add some basic firewall rules to our VM. The course repository has a set of predefined basic firewall rules that we can use. Let's see what they are:

```
1 ubuntu@wse380:~$ cd wse380-webapp
2 ubuntu@wse380:~/wse380-webapp$ cat setup/ufw/install-rules-basic.sh
3 #!/bin/bash
4
5 # default rules
6 sudo ufw default deny incoming
7 sudo ufw default allow outgoing
8
9 # http rules
10 # sudo ufw allow in 80/tcp
11 # sudo ufw allow in 443/tcp
12
13 sudo ufw allow in 22/tcp
14
15 sudo ufw enable
16 sudo ufw status verbose
17 ubuntu@wse380:~/wse380-webapp$
```

As you may have noticed, this file is a *.sh* file, which indicates that it is a *shell script*. These shell scripts contain commands that we can execute in the command line. In particular, this script is a *bash script* (bash is a specific shell), which we can tell from the *shebang* in the first line (*#!/bin/bash*). We can execute this shell script (which executes all the commands sequentially) with the following:

```
1 ubuntu@wse380:~/wse380-webapp$ ./setup/ufw/install-rules-basic.sh
2 Default incoming policy changed to 'deny'
3 (be sure to update your rules accordingly)
4 Default outgoing policy changed to 'allow'
5 (be sure to update your rules accordingly)
6 Rules updated
7 Rules updated (v6)
8 Command may disrupt existing ssh connections. Proceed with operation (y|n)? y
9 Firewall is active and enabled on system startup
10 Status: active
11 Logging: on (low)
12 Default: deny (incoming), allow (outgoing), disabled (routed)
13 New profiles: skip
14
```

15	To	Action	From
16	--	-----	----
17	22/tcp	ALLOW IN	Anywhere
18	22/tcp (v6)	ALLOW IN	Anywhere (v6)

2.2.3 Fail2Ban [4]

Fail2Ban is a traditional intrusion-prevention software that can run on Linux systems. In this course, we will be using fail2ban to add a layer of protection against attackers who may try to compromise our server. Let's install it by running:

```
1 sudo apt update
2 sudo apt upgrade
3 sudo apt install fail2ban
```

Similarly to ufw, let us add some basic configuration rules. The course repository has a set of predefined basic rules that we can use:

```
1 ubuntu@wse380:~/wse380-webapp$ cat setup/fail2ban/fail2ban.conf
2 [DEFAULT]
3 bantime = 1d
4 findtime = 1d
5 ignoreip = 127.0.0.1/8 192.168.0.0/16 130.245.0.0/17
6 maxretry = 3
7
8 banaction = ufw
9 banaction_allports = ufw
10
11 [sshd]
12 enabled = true
13
14 [ufw]
15 enabled = true
16 filter = ufw
17 logpath = /var/log/ufw.log
18 ubuntu@wse380:~/wse380-webapp$
```

We can see that there is a default policy that bans IP addresses for one day if they trigger 3 failed attempts in one day, excluding local IP addresses (127.0.0.1/8 and 192.168.0.0./16) and Stony Brook IP addresses (130.245.0.0./17). Afterwards, the configuration tells fail2ban to check for SSH and UFW logs, among things.

We can install these rules by running the following script:

```
1 ubuntu@wse380:~/wse380-webapp$ ./setup/fail2ban/copy-config.sh
2 Waiting 5s for fail2ban to restart...
3 Status
4 |- Number of jail: 1
5 '- Jail list: sshd
6 ubuntu@wse380:~/wse380-webapp$
```

Now we have finished the basic access control configuration for our servers, and can begin installing the main prerequisite we will need for the course: Docker.

2.3 Installing Docker and Docker Compose

Docker is a tool that allows developers to package their programs into portable containers. For the sake of this course, all you need to know is that Docker removes all of the complexity in setting up an application. We will not go into any more detail on Docker (aside from a few basic commands), although I encourage you to continue reading up on Docker as it is an industry standard in the deployment of web applications.

To install Docker, we are going to use the apt tool we discussed earlier. We first want to update apt to make sure it downloads the latest version of each application we request. We do that using the following commands in `docker/install_docker.sh` [2]:

```

1 ubuntu@wse380:~/wse380-webapp$ cat setup/docker/install_docker.sh
2 #!/bin/bash
3
4 # Add Docker's official GPG key:
5 sudo apt-get update
6 sudo apt-get install ca-certificates curl
7 sudo install -m 0755 -d /etc/apt/keyrings
8 sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg -o /etc/apt/keyrings/docker.asc
9 sudo chmod a+r /etc/apt/keyrings/docker.asc
10
11 # Add the repository to Apt sources:
12 echo "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.asc] \
13     https://download.docker.com/linux/ubuntu \
14     $(. /etc/os-release && echo "$VERSION_CODENAME") stable" | \
15     sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
16 sudo apt-get update
17
18 # Install
19 sudo apt-get install docker-ce docker-ce-cli containerd.io \
20     docker-buildx-plugin docker-compose-plugin

```

We have now successfully installed Docker! One more thing — we will also let our user account to use Docker without root privileges by running the following:

```

1 # Use docker without sudo
2 sudo usermod -aG docker ${USER}
3 su - ${USER}

```

This will add the user account into the ‘docker’ user group for the correct permissions, and relogin to refresh the groups. Alternatively, we can also close the current SSH connection and reconnect for the groups to take effect.

Verify your installation by running:

```

1 docker run hello-world

```

Now let us install Docker Compose to help with running multiple Docker containers [1].

```

1 sudo apt-get update
2 sudo apt-get install docker-compose-plugin

```

Verify the Docker Compose installation by running:

```

1 ubuntu@wse380:~/wse380-webapp$ docker compose version
2 Docker Compose version v2.24.6

```

Next Steps

In the next session, we will setup our web application and register a domain name. Start to think about what domain name you would like to use! Please note that the domain name will be restricted to a low-cost TLD such as .shop, .online, and .store.

3 Session 2: Web Application Setup

Today, we are going to walk through a hands-on guide for installing the software we need for our servers.

3.1 Web Application Infrastructure

Servers are machines that are designed to be running software 24/7. However, there are different ways to deploy software on machines: directly on the *bare metal*, or alternatively with *virtualization* (i.e., by using a *virtual machine* (VM) or a *container*). There are differences in these methods that bring certain (dis)advantages, and one method may be better than the others depending on the scenario (usually for performance reasons). In general, virtualization techniques are supposed to be able to provide a common ground by defining standard environments for applications.

In this class, we will not elaborate on these differences, but it is important to recognize that software is now typically deployed using virtualization. You will be given access to a *virtual machine*, and we will be running software directly in the operating system of the VM, and inside Docker containers in the VM.

3.2 Docker Compose Services

Let's begin configuring the web application software by using Docker Compose. Compose is a plugin that makes it easy to manage multiple Docker containers in one context. Typically, one Compose project is created for one application, which may comprise multiple smaller pieces of software. In this class, we will be using a provided Compose file to begin configuring our individual projects.

3.2.1 Environment Variables

Environment variables, as their name suggests, are variables that exist in the environment of an application, and may not actually exist in the application configuration itself. For example, something analogous to an environment variable might be whether there is still sunlight outside. Depending on whether there is still light outside, you may opt to go to a restaurant, or go to a late-night drive-through for a fast food chain.

Applications can read values from their environment in a similar way. Let's change directory into our working directory named “~/wse380-webapp/wordpress-compose”, and begin to configure the environment variables used by our software [3].

```
1 cd ~/wse380-webapp/wordpress-compose
2 nano .env
```

This should have opened the `.env` file in the nano text editor and we should see something like this:

```
1 MYSQL_DATABASE_NAME=wordpress
2 MYSQL_ROOT_PASSWORD=root_password
3 MYSQL_USER=wordpress_database_user
4 MYSQL_PASSWORD=wordpress_database_password
5
6 IP=yourip
7 DOMAIN=yourdomain
8 EMAIL=youremail
```

Modify the values for the following variables such that it remains in the format `VARIABLE=value`, with one per line (note that the values are case sensitive):

- `MYSQL_ROOT_PASSWORD` - the root password for your database; please use a secure password,
- `MYSQL_PASSWORD` - the password of the database account used by WordPress; please use a *different* secure password,
- `MYSQL_USER` - the username of the database account used by WordPress; please select any username,

- IP - the IP address of your VM; please input the IP address here,
- DOMAIN - the domain name you will register; please leave it be for now,
- EMAIL - the email you want to associate with TLS certificate expiry reminders; please leave it be for now.

Note that this file now contains *secrets for your application that should not be shared*. In deployment environments, it is important to make sure these files are not accidentally indexed or stored (e.g., by version control software such as git).

Next, please run the following to update the configuration files:

```
1 ubuntu@wse380:~/wse380-webapp$ ./wordpress-compose/scripts/replace-ip.sh
```

Now, let's start our applications!

```
1 ubuntu@wse380:~/wse380-webapp$ cd wordpress-compose
2 ubuntu@wse380:~/wse380-webapp/wordpress-compose$ docker compose up -d
3 ...
4 [+] Running 3/3
5  ✓ Container db           Running      0.0s
6  ✓ Container wordpress Running      0.0s
7  ✓ Container webserver Running      0.0s
```

Your VM will begin building the application containers and running them. After a short period of time, you should see similar text at the end, saying that all 3 containers db, wordpress, and webserver are running.

3.3 WordPress Installation via the Web Interface

One of the typical features of web apps (particularly for large CMS software) is the ability to install the application via the web. Let's see what it looks like: go to your browser and enter the URL `<your_ip>:12345`. You should now see the WordPress installation page!

Follow these steps to complete the installation:

1. Select a language (e.g., English)
2. Enter a site title (e.g., WSE380 - John's Blog)
3. Enter a username (e.g., admin-username)
4. Enter a secure password (or keep the default generated one) and store it in a password manager (e.g., your browser's)
5. Enter your email
6. Press the "Install WordPress" button!

The page should change after a short period of time, informing you of a successful installation and to log in with the username and password that was just provided. Please do so – we should land on the URL `http://<your_ip>/wp-admin/index.php` with a big "Welcome to WordPress!" banner.

3.4 DNS

So far we have accessed our web application directly via its IP address, but we want to be able to associate a memorable, human-readable name for our site! Follow along and let's proceed to register our own domain name!

3.4.1 Domain Registration

We register domain names through *domain registrars*, such as Dynadot, Namecheap, and GoDaddy. These registrars enable us to search for a domain name and tell us whether they are already registered, or available for purchase. Search for the domain you want to register!

3.4.2 DNS Name Servers

After we purchase (more accurately, lease!) a domain, we can configure its *authoritative name server*. This name server is a server that we identify by its IP address, and is responsible for translating the domain into an IP address. We typically configure the authoritative name server directly via the domain registrar from which we purchased the domain. Using the registrar, configure a DNS A record that points to your IP address.

3.4.3 Updating the Nginx/WordPress Configuration

Now we need to update our configurations to include the domain.

1. First, update our environment variable for the DOMAIN variable by running:

```
1 ubuntu@wse380:~/wse380-webapp$ nano wordpress-compose/.env
```

and then modifying the domain value to your newly registered domain.

2. Run the script

```
1 ubuntu@wse380:~/wse380-webapp$ ./wordpress-compose/scripts/replace-domain.sh
```

3. Next, we need to tell WordPress that it should respond to requests for the domain that we just purchased! Head to Settings in the admin panel, which should open the General Settings page. Update the following:
 - WordPress Address (URL) to `http://<your_domain>`
 - Site Address (URL) to `http://<your_domain>`
4. Afterwards, scroll down and hit the Save Changes button.
5. Verify that you can reach your WordPress site by navigating to `http://<your_domain>` in your browser! You will need to re-authenticate to the admin panel, since the origin is now different.

3.5 TLS Certificates

Now, we have established a site we can access via a domain name. Our next step is to obtain a TLS certificate — the missing piece that will enable our browsers to make secure connections via HTTPS (the lock icon in the URL bar)! The entire ecosystem is complicated, but the key pieces are:

- we will request a certificate from a globally trusted Certificate Authority (CA),
- the CA will verify that we own the domain, and
- the CA will issue a certificate that says we own the specified domain.

After we obtain the certificate, we can update our configuration so that our Nginx web server knows to present the certificate for browsers that request HTTPS connections to our website. The browser will see that our valid certificate is from a globally trusted CA, and thus trust that we are the actual owner of the domain.

3.5.1 Certificate Request

Let's begin the process of obtaining the certificate.

1. Edit the `wordpress-compose/docker-compose.yml` file with nano, and remove the leading # symbols in front of the certbot service to enable it.

2. Restart the containers by running

```
1 ubuntu@wse380:~/wse380-webapp/wordpress-compose$ docker compose up -d
```

3. Monitor the certificate process by running:

```

1 ubuntu@wse380:~/wse380-webapp/wordpress-compose$ docker compose logs -f certbot
2 ...
3 certbot | Requesting a certificate for <your_domain> and www.<your_domain>
4 certbot |
5 certbot | Successfully received certificate.
6 certbot | Certificate is saved at: /etc/letsencrypt/live/<your_domain>/fullchain.pem
7 certbot | Key is saved at: /etc/letsencrypt/live/<your_domain>/privkey.pem
8 certbot | This certificate expires on 2024-06-01.
9 certbot | These files will be updated when the certificate renews.
10 certbot | NEXT STEPS:
11 certbot | - The certificate will need to be renewed before it expires. Certbot can
   automatically renew the certificate in the background, but you may need to take steps to
   enable that functionality. See https://certbot.org/renewal-setup for instructions.

```

4. Verify that you can see the certificates:

```

1 ubuntu@wse380:~/wse380-webapp/wordpress-compose$ docker compose exec webserver ls -la /etc/
   letsencrypt/live
2 total 16
3 drwx----- 3 root    root      4096 Mar 3 21:59 .
4 drwxr-xr-x  7 root    root      4096 Mar 3 22:18 ..
5 -rw-r--r--  1 root    root        740 Mar 3 21:59 README
6 drwxr-xr-x  2 root    root      4096 Mar 3 22:18 <your_domain>

```

5. Now, re-comment the lines in wordpress-compose/docker-compose.yml for the certbot service (by adding the leading # symbols) so we do not send unnecessary certificate requests.

3.5.2 Updating Docker Compose

Now, the final step is to tell Docker Compose to use this new certificate. Let us run the following:

```

1 ubuntu@wse380:~/wse380-webapp$ ./wordpress-compose/scripts/change-to-https.sh
2 ubuntu@wse380:~/wse380-webapp$ cd wordpress-compose
3 ubuntu@wse380:~/wse380-webapp/wordpress-compose$ docker compose up -d

```

Verify that you are now able to use a secure connection to your website by specifying HTTPS in your browser!

Next Steps

In the next few sessions, we will be focusing on customizing our sites. Try to come up with a few ideas for what you would like your site to focus on! These domains that you registered today are yours to keep — just keep in mind that they will expire in one year from today!

4 Session 3: WordPress Customization

WordPress provides a content management system, and is one of the most popular of those publicly available. In this lesson, we will be learning how to interact with the basic functionalities offered by our WordPress site. For each section below, we will start from the administrative panel of our WordPress sites (located at https://<your_domain>/wp-admin).

4.1 Content

One of the most important reasons why we would want to have an online site is to post our own content! Let's get started by adding some content to our empty site. Feel free to substitute the example content below with any content you desire (just remember that all published content will be public!).

4.1.1 Categories

Let's create a few categories to organize our content.

1. Go to "Posts → Categories"
2. In the form for "Add New Category", create the categories with the following names, slugs, and descriptions:

Name	Slug	Description
Animals	animals	Category for animal posts
Cooking	cooking	Category for cooking posts
Anything else you want!	Anything else you want!	Anything else you want!

4.1.2 Tags

Let's create some tags to further organize our content (similar to categories, but usually not topic related).

1. Go to "Posts → Tags"
2. In the form for "Add New Tag", create the tags with the following names, slugs, and descriptions:

Name	Slug	Description
Images	images	Tag for posts with images
Tutorials	tutorials	Tag for posts with tutorials
Anything else you want!	Anything else you want!	Anything else you want!

4.1.3 Posts

Now that we have some organizational markers, let's make new posts.

Post 1: Animal Pictures. Let's make a post to share some of our favorite animal pictures! Follow along on screen or below:

1. Go to "Posts → Add New Post", which should redirect your page to a "what you see is what you get" (WYSIWYG) content editor.
2. Use a descriptive title (e.g., "Fluffy Dog Pictures")

3. Write a description for the post (e.g., “I would like to share some of the floofiest dogs I have ever seen!”)
4. Hit “Enter” to start a new *block*
5. Type “/image” and select the “Image” suggestion
6. Upload an image!
7. Create a descriptive heading before the image by typing in text, clicking the paragraph symbol, and selecting “Transform to Heading”
8. Click the “Settings’ menu bar icon on the right of “Publish”
9. Scroll down to the “Categories” section, toggle the “Animals” category, and then unselect the “Uncategorized” category
10. Click “Save draft” at the top right
11. Click the “Preview” laptop icon on the left of the “Publish” button to preview what your post will look like on different device screen sizes
12. Click “Preview in new tab” under the “Preview” laptop icon

Post 2: Food Pictures. Let’s make a new post to share some of our favorite foods! Follow along on screen or below:

1. From the preview screen of the last post, you should be able to click “+ New → Post” on the top menu to create a new post
2. Use a descriptive title (e.g., “Best Food Pictures”)
3. Write a description for the post (e.g., “I would like to share some of the best foods I have ever had!”)
4. Hit “Enter” to start a new *block*
5. Type “/gallery” and select the “Gallery” suggestion
6. Upload several images this time!
7. Click the “Settings’ menu bar icon on the right of “Publish”
8. Scroll down to the “Categories” section, toggle the “Cooking” category, and then unselect the “Uncategorized” category
9. Click “Save draft” at the top right
10. Click the “Preview” laptop icon on the left of the “Publish” button to preview what your post will look like on different device screen sizes
11. Click “Preview in new tab” under the “Preview” laptop icon

Tutorial: How to Cook (your favorite dish).

1. From the preview screen of the last post, you should be able to click “+ New → Post” on the top menu to create a new post
2. Use a descriptive title (e.g., “How to Cook Rice”)
3. Write a description for the post (e.g., “I would like to share a recipe for the most delicious rice I have ever had!”)
4. Create some sections for your post (e.g., “Ingredients Needed”, “Steps”)
5. Hit “Enter” to start a new *block*
6. Type “/list” and select the “List” suggestion
7. Click the “Select List” button on the left of the toolbar to switch between a numbered list and a bulleted list

8. Fill in the recipe!
9. Add an image of the final dish
10. Click the “Settings” menu bar icon on the right of “Publish”
11. Scroll down to the “Categories” section, toggle the “Cooking” category, and then unselect the “Uncategorized” category
12. Scroll down to the “Tags” section, and enter “Images” and “Tutorials” to add the tags we created before
13. Click “Save draft” at the top right

4.2 Styling and Themes

Now that we have some content on our site, we might be interested in changing how our site looks. Follow along on screen or below:

1. Go to “Appearance → Themes”
2. Click on “Add New Theme”
3. Browse the directory until you find a theme that you would like to install
4. You can preview themes before choosing to install with the “Preview” button
5. When you have found a theme you like, press “Install”, and then “Activate”
6. Go back to “Posts → All Posts” to preview how the posts look in the new theme

Repeat as necessary until you find a theme that you are satisfied with!

4.3 Plugins

Plugins are a popular method to distribute third-party additions to existing software. There is no strict definition, as the method for writing and loading plugins differs from one software to another, but, in general, they typically offer advanced functionality that is not present within the vanilla software itself. For example, plugins may offer features such as security and privacy, performance optimizations, or analytics and tracking. We will conclude this session by adding several plugins to our WordPress sites.

4.3.1 Security - Wordfence

Whenever we maintain a site on the Internet, we open up our site to both benign and malicious traffic. We should be particularly concerned because our new site is a WordPress site – WordPress is one of the largest content management system software in the market, and attackers routinely scan the web to search for vulnerabilities in WordPress sites. Thus, the very first plugin we will be installing is a security plugin. Follow the steps below:

1. Go to “Plugins → Add New Plugin”
2. Search for “Wordfence”
3. Install the plugin titled “Wordfence Security — Firewall, Malware Scan, and Login Security”
4. Go to “Plugins → Installed Plugins”
5. Press “Activate” under the Wordfence plugin that we just installed, and follow the steps
 - (a) The Wordfence licensing page should open
 - (b) Subscribe for a free license with your email
 - (c) Check your inbox for the license key
 - (d) Activate your key manually or visit the link in the button in the email

Now, let us check what we have just installed by examining the new Wordfence section on our admin dashboard.

Firewall. A firewall is an essential component of applications. Similar to the `ufw` we installed for our servers, Wordfence comes with its own firewall. The main difference is that this firewall is a web application firewall (WAF) — it is designed to protect against malicious payloads specifically for WordPress.

Let us configure one setting in our firewall.

1. Scroll down to and open the “Rate Limiting” section
2. For the “If anyone’s requests exceed” option, select “120 per minute” and “throttle it”
3. Click “SAVE CHANGES” at the top

Now, the Wordfence WAF will automatically throttle (slow down) the responses to individual IP addresses who send over 120 requests per minute.

Scanner. Another essential security component is a server-side scanner. These scanners typically scan files on the server to see if they detect any malicious snippets of code. These scanners are highly useful, particularly for CMS sites, as there are generic, core files that are used across all deployments that should have predictable content. One of the main tasks of these scanners is to ensure that these core files are not malicious.

Login Security. As you may have observed with many modern web applications, a traditional username and password login combination are not very secure, for a number of reasons. Wordfence comes with a built-in two-factor authentication (2FA) support. If you have a 2FA time-based one-time password (TOTP) application (e.g., Google Authenticator), you can proceed to configure 2FA for your WordPress admin user by following the steps.

Tools. Let us turn on traffic logging for all traffic by going to “Live Traffic Options” and selecting “Traffic logging mode: ALL TRAFFIC”. Now, we should be able to see traffic logs in real time.

4.4 Free Time: Site Customization

Now let’s work on customizing our sites! As a start, we’ll work on brainstorming ideas of what we can add, and then you will have the rest of this session as free time to work on your sites. Feel free to further explore the admin panel to add content, change the styling, look into more plugins, etc. As always, if you have any questions, we can work through them together.

The goal of this customization time is for you to prepare your site for real visitors, and for us to analyze the traces of their interaction with your site in the next sessions. I encourage you to take some time to customize your sites, and invite your friends, classmates, family, etc. to visit it! At the end of this rotation, we will get a chance to learn about what everyone else has created, and compare visitor statistics.

5 Session 4: Website/Server Monitoring

One of the key classes of software necessary for the success of web applications is monitoring software. This type of software is dedicated to providing essential metrics such as whether your website is responding and if it is overloaded in t

5.1 Resource Usage

Servers are machines that have finite resources. Resources that are primarily of interest when deploying applications are the trifecta of *CPU*, *memory*, and *disk*. Central processing units (CPUs) process every task on a machine; random access memory (RAM) enables applications to track their required data without having to always read from and write to the actual disk; disk refers to the piece of hardware that stores data for long periods. If these resources approach their limits, the applications running on that machine will be severely affected.

For this reason, we will be installing a plugin to monitor the resources used by the WordPress server.

1. Go to “Plugins → Add New Plugin”
2. Search for “Server Health Stats”
3. Install the plugin titled “WP Server Health Stats”
4. Go to “Plugins → Installed Plugins”
5. Press “Activate” under the WP Server Health Stats plugin that we just installed, and follow any installation steps

Now you should see server resource usage stats at the bottom of every page, and in the main admin dashboard!

5.2 Uptime

Another piece of essential monitoring software is an uptime monitor that regularly checks to see if your website is responsive, and alerting you if it is not responsive (or responding very slowly). We can use a simple free service to do this for us. Follow along on screen or below:

1. Go to <https://uptimerobot.com>
2. Create a free account
3. Click the “New monitor” button
4. Enter your site address under “URL to monitor” (e.g., https://<your_domain>)
5. Leave the other options as their default values and click “Create monitor” at the bottom

You should be able to see that the UptimeRobot service successfully confirmed that your website is up!

5.3 More Plugins

Here we list several ideas for additional plugins that you can install on your site. We will walk through the installation of the first, Google Site Kit, and leave the rest as optional plugins. As always, feel free to install other plugins to your site, but maintain a cautious customer mindset — the average rating, and number of reviews, is a proxy metric for the quality of a plugin.

5.3.1 Analytics - Google Site Kit

As a site administrator, it is difficult to get insight into visitor-related metrics (e.g., number of visitors) without using analytics software. One of the more popular ones is Google Analytics, and the corresponding WordPress plugin is titled “Site Kit by Google”. Let’s install the plugin by:

1. Go to “Plugins → Add New Plugin”
2. Search for “Site Kit by Google”
3. Install the plugin titled “Site Kit by Google — Analytics, Search Console, AdSense, Speed”
4. Go to “Plugins → Installed Plugins”
5. Press “Activate” under the Google Site Kit plugin that we just installed, and follow the steps on screen

5.3.2 Performance - LiteSpeed Cache (Optional)

1. Go to “Plugins → Add New Plugin”
2. Search for “LiteSpeed Cache”
3. Install the plugin titled “LiteSpeed Cache”
4. Go to “Plugins → Installed Plugins”
5. Press “Activate” under the LiteSpeed Cache plugin that we just installed, and follow any installation steps

5.3.3 Forms - WPForms (Optional)

1. Go to “Plugins → Add New Plugin”
2. Search for “WPForms”
3. Install the plugin titled “Contact Form by WPForms – Drag & Drop Form Builder for WordPress”
4. Go to “Plugins → Installed Plugins”
5. Press “Activate” under the WPForms plugin that we just installed, and follow any installation steps
6. Click the “Create Your First Form” button on the onboarding page
7. Add a “Simple Contact Form” to a new page called “Contact Us”

5.3.4 Forms - WP Mail SMTP (Optional)

Currently, the free version of the WPForms plugin we are using does not support storing the submitted form information. The only delivery mechanism it supports is email (sending an email when a form is submitted). To setup the email notification system, we will need to do the somewhat complicated procedure below:

1. Go to “Plugins → Add New Plugin”
2. Search for “WP Mail SMTP”
3. Install the plugin titled “WP Mail SMTP”
4. Go to “Plugins → Installed Plugins”
5. Press “Activate” under the WP Mail SMTP plugin that we just installed, and follow any installation steps
6. Setup integration with one of the mailers (e.g., Brevo) to receive email notifications
 - This will require you to have an account with the third party SMTP server
 - In addition, you will need to go through various processes to obtain an API key
 - Finally, you will most likely have to setup additional DNS records for your domain

5.4 Free Time: Site Customization

Now let's work on customizing our sites again! Similar to last time, we'll work on brainstorming ideas of what we can add, and then you will have the rest of this session as free time to work on your sites. Feel free to further explore the admin panel to add content, change the styling, look into more plugins, etc. As always, if you have any questions, we can work through them together.

Next Steps

Our next session will be the last one for the rotation! We will go through everyone's site to display what everyone has worked on the past few weeks, and learn how to identify visitor behavior. Before our next class, please try to:

- Create a post that summarizes what you have learned from this rotation, and any further ideas you would like to have explored. If you would prefer to leave feedback anonymously, we will also be sending out a Google Form to collect anonymous responses!
- Clean up and personalize your site by removing the general placeholder text.
- Try to have other people visit your site — the more the better!

6 Session 5: Server Traces & Site Showcase & Review

As a site admin, we need to be able to determine how our visitors are interacting with our application. Furthermore, we need to be able to do this from various perspectives — e.g., from the perspective of our server, our firewalls, and our applications. Every piece of software typically has their own mechanism that produce log messages with which we can attempt to piece together an understanding of what previously happened.

6.1 Understanding Logs

6.1.1 Server Logs

One of the places where we can find log messages is directly on the server. In our case, we can find most of the log messages stored inside the directory `/var/log`. Let's take a look at the server log files.

First, inspect what log files are available by running:

```
1 ubuntu@wse380:~/wse380-webapp$ ls -lh /var/log
2 total 75M
3 ...
4 -rw-r----- 1 syslog adm      113K Mar 28 19:48 auth.log
5 -rw-r----- 1 syslog adm      2.0M Mar 23 23:41 auth.log.1
6 -rw-r----- 1 syslog adm      422K Mar 16 23:59 auth.log.2.gz
7 -rw-r----- 1 syslog adm      246K Mar 10 00:00 auth.log.3.gz
8 -rw-r----- 1 syslog adm      1.7M Mar 3 00:00 auth.log.4.gz
9 ...
10 -rw-r----- 1 root   adm      110 Mar 24 00:00 fail2ban.log
11 -rw-r----- 1 root   adm     220K Mar 23 23:41 fail2ban.log.1
12 -rw-r----- 1 root   adm      57K Mar 16 23:40 fail2ban.log.2.gz
13 -rw-r----- 1 root   adm      60K Mar 9 23:58 fail2ban.log.3.gz
14 ...
```

Let's take a look at the following log files:

- **`/var/log/fail2ban.log`**. Run the following command:

```
1 ubuntu@wse380:~/wse380-webapp$ less /var/log/fail2ban.log
```

What can you tell me about this log file?

- each line represents one event,
- an event details a connection attempt or punitive action (e.g., a ban)

- **`/var/log/auth.log`**. Run the following command:

```
1 ubuntu@wse380:~/wse380-webapp$ less /var/log/auth.log
```

What can you tell me about this log file?

- each line represents one event,
- events consist of session openings and closings, connection details, etc.

6.1.2 Application Logs.

Another place where log messages are typically gathered are in applications. Each application has their own logging mechanism and destination, so there is no universal place to easily check for logs (`/var/log` on Linux is the closest there is to a universal directory). Let's inspect the logs for our web server and WordPress.

nginx. Let's take a look at what the nginx web server log files look like. Run the following commands:

```
1 ubuntu@wse380:~/wse380-webapp$ wordpress-compose/goaccess/copy_log.sh
2 ubuntu@wse380:~/wse380-webapp$ less wordpress-compose/goaccess/access.log
```

What can you tell me about this file?

- each line represents one request,
- the log message includes data such as the IP address of the user, the date and time of the request, the domain name, etc.

Run the following commands:

```
1 ubuntu@wse380:~/wse380-webapp$ setup/goaccess/install_goaccess.sh
2 .....
3 ... this should prompt you to install several packages, accept them ...
4 ... then it will install and compile a new program called goaccess ...
5 .....
6 ubuntu@wse380:~/wse380-webapp$ wordpress-compose/goaccess/run_goaccess.sh
```

WordPress/Wordfence. Now, let us take a look at the logs inside Wordfence itself. Follow along on screen or below:

1. Go to https://<your_domain>/wp-admin
2. Go to “Wordfence → Tools”
3. Inspect the Live Traffic log messages displayed

6.2 Site Showcase

Now, let us wrap up our rotation! We will take turns to showcase the sites we built in the past few weeks.

6.3 Conceptual Review

In this WSE380 rotation, we covered the technical foundations of web applications. In particular, we accomplished the following tasks:

- familiarity with the Linux command line
- DNS
 - domain name registration and DNS record configuration,
 - TLS certificate requests,
- WordPress
 - deploying our self-hosted WordPress application
 - customizing and personalizing our own sites
- installing various security software (e.g., ufw, fail2ban, and Wordfence), other web admin software, and plugins
- understanding log messages, which are traces that are produced from interacting with the server

Thank you all for signing up for this rotation! We will be reaching out to send an anonymous Google Form to collect feedback for this rotation. Again, please feel free to email me with any further questions you may have, and good luck with the rest of your semesters!

References

- [1] Install the compose plugin. <https://docs.docker.com/compose/install/linux/>.
- [2] Install docker engine on ubuntu. <https://docs.docker.com/engine/install/ubuntu/>.
- [3] Kathleen Juell. How to install wordpress with docker compose. <https://www.digitalocean.com/community/tutorials/how-to-install-wordpress-with-docker-compose>, 2024.
- [4] Jean-Jerome Levy. Turn your nginx server into a fortress with fail2ban and ufw. <https://scalastic.io/en/ufw-fail2ban-nginx/>, 2023.